

Cloud Operations, FS2022

Luzia Kündig

July 11, 2022

1 Introduction

CI/CD	Deployment environment, provides name and url for monitoring
Ansible	Installation and configuration of applications
Terraform	Provisioning of infrastructure
Kubernetes	Configuration + application: Deployment, scaling, managing workloads
Helm	Package- and Lifecycle Manager for K8s
Kustomize	Overlaying declarative specifications on top of existing K8s Manifests
Prometheus	Monitoring K8s Infrastructure and applications for reliability
Service Mesh	Traffic Management, Security, Observability and Service Discovery
GitOps	Everything as code, declarative system operation definition, control loop

1.1 DevOps




Figure 1: DevOps Cycle

The DevOps pipeline focuses on Continuous Integration / Continuous Deployment. It applies a systems thinking and avoids too much focus on only one piece of the workflow.

It values feedback, automated and personal to keep production healthy. Monitoring is an important type of automated feedback.

Learning and continued experimentation lead to constantly improved systems and workflows.

1.1.1 Plan

Add objectives and requirements to the **backlog**, feedback from end users and operations team. Add backlog **tasks to sprints**. Track and plan activities using **board and project management tools**.

1.1.2 Code

Code from all developers gets integrated into a **central source code repository**.

1.1.3 Build

Continuous Integration pipeline is invoked every time code is pushed into the central repository. Includes **automated unit and integration tests**. Only after successful build and test, code can be reviewed and merged.

1.1.4 Test

Automated deployment of code into a testing environment. Tests executed include **load, accessibility, performance and end-to-end testing**. Manual work like **user acceptance testing** also usually happens at this stage.

1.1.5 Release

Tag a snapshot of the code with a **semantic versioning number**. Changes, features, breaking changes and deprecated features are documented. Release can include artifacts such as **binaries** and **packages**.

1.1.6 Deploy

Installs release into **production environment**. Can be automated or manual.

1.1.7 Operate

Infrastructure and Operations team ensure **smooth operation of the product**. **Scaling infrastructure** to meet demands.

Issues in infrastructure can be troubleshooted and resolved. **Document issues** for next planning stages.

1.1.8 Monitor

Collect data on usage, performance, errors and more. Data collected is used for next iteration of DevOps cycle.

1.2 Automated DevOps with GitLab CI

Using GitLab CI pipeline, code can be built, tested and deployed automatically on every change. Tasks are defined in `.gitlab-ci.yml` file kept together with the code repository.

- 1.2.1 Integrating Kubernetes
- 1.2.2 Auto-DevOps: Setup with zero configuration
- 1.2.3 Creating gitlab CI Configuration
- 1.2.4 Testing with Docker

2 Service Mesh

2.1 Microservice Principles

The modern architecture of microservices focuses on the following aspects of operation.

1. Deployment Independence
2. Organized by business capability
3. Products not Projects
4. API Focused
5. Smart endpoints and dumb pipes
6. Decentralized governance
7. Decentralized data management
8. Infrastructure Automation (IaC)
9. Design for failure
10. Evolutionary design




Figure 2: Traditional Service Architecture

2.2 Service Mesh

A service mesh is a dedicated and configurable infrastructure layer that handles the communication between services without having to change the code in a microservice architecture.

Instead of every service implementing important functionality themselves, a proxy is deployed to intercept network traffic to each container inside a pod.




Figure 3: Service Mesh Architecture

Service Mesh is responsible for Traffic Management, Security, Observability and Service Discovery inside each microservice. This is done by embedding capabilities via sidecar containers directly into kubernetes pods.

2.3 The Istio Approach

Grafana	Visualization of metrics collected by Prometheus
Istio Ingress Gateway	Envoy Proxy which serves as an entry point for the service mesh
Istiod	Backbone of the Istio control plane, configures sidecar proxies
Jaeger	Provides tracing functionality for traffic inside the service mesh
Kiali	Istio dashboard
Prometheus	Collects monitoring information from sidecar proxies

2.4 Canary Deployment

Canary deployments are used to test a new version of a microservice in production. To do so, the new version gets deployed, but only a small percentage of traffic gets sent to it. If no problems occur and no customers complain, the traffic rate to the new version gets increased.

2.5 Virtual Services

Inside a VirtualService configuration file, a set of traffic routing rules can be defined that should be applied when a host is addressed. Each routing rule defines matching criteria for traffic of a specific protocol.

If traffic is matched, it is sent to a named destination service or a specific subset / version of it. This can be used for load balancing and different deployment strategies.

2.6 Jaeger Spans and Traces

A trace represents a single request through the service mesh which gets handled by the services. Each unit of work inside a trace is called a span. Spans can be requests to other services, for example.

3 GitOps

The concept of GitOps enforces using Git as *single source of truth defining the application state*.

Instead of different people using *kubectl create / apply* or *helm install / upgrade* commands from their own laptops, the whole configuration of a kubernetes cluster is kept inside a (separate!) Git repository.

This brings the known advantages of a version control system to Continuous Delivery for Kubernetes Clusters:

- **Versioned and immutable**
- **Declarative** definitions of apps, environments and configurations
- **Automated and repeatable**, less opportunity for errors
- Code review for changes
- Tracking of who did what
- Rollback via Git
- Whole infrastructure can be recreated from source control

Separating the git repository holding kubernetes manifests and the actual application code provides more simplicity when updating some deployment information.

No build and test pipeline will be triggered without any change to the actual code. Git history will be cleaner. Application may be distributed over several git repositories. Separation of access is possible.

3.1 ArgoCD


Declarative GitOps tool used to deploy applications on Kubernetes.

It is lightweight, easy to configure, built with GitOps Workflow in mind and intended only for Kubernetes.

How does it work?

- Spins up its own controller inside the cluster
- watches for changes in a repository
- compares against resources deployed in the cluster
- synchronizes both states (desired state wins)

Its key concepts are the following:



App/Application	Group of K8s resources defined by a manifest (CRD), used to monitor repository and update cluster
Application source type	Which build tool is used to build the app
Live State	Current state of the application cluster
Target State	Desired state as stored in Git
Sync Status	Is live the same as target?
Sync	Updating app to the target state
Sync operation status	Success/failure state of the sync operation
Refresh	Comparison of Code in Git with live state (.. of Sync Status)
Health Status	Is the app running correctly? Serving Requests?
Configuration Management Tool	Tool to create manifests from files in direcorey (Kustomize, Ksonnet)
Configuration Management Plugin	Custom tools ...

3.2 ArgoCD Application

Specifies information such as ArgoCD project, source repo, revision, path, cluster, namespace.

Can be created via command line, Web Interface, Yaml in web, K8s Manifest (CRD).

Can be specified using kustomize, helm, ksonnet, jsonnet, plain yaml, custom configuration management tool set up as plugin to ArgoCD.

```
argocd app create, argocd list
```

Different app health statuses include

- Progressing
- healthy
- Degraded
- Suspended
- Missing

3.2.1 History and Rollback

ArgoCD keeps track of the various versions deployed and allows you to go back to a previous state. Can be accessed from the graphical user interface.

3.2.2 Manual Sync

- **Prune:** Allow deleting resources that are unexpected, meaning they no longer exist in git.
- **Dry Run:** Preview what an apply operation would do without affecting the cluster
- **Apply only:** Skip pre/post sync hooks
- **Force:** Deletes and re-creates resource(s) when patch encounters conflict after 5 retries

3.2.3 Automated Sync

Only occurs if App Sync Status is *OutOfSync*. Attempts one sync per unique combination of commit hash and parameters of the app – unless *selfHeal* flag is true.

selfHeal attempts to sync after a default timeout of 5 seconds.

3.2.4 ArgoCD Projects

Projects can be a logical group of applications, which supports organization by teams.

Features include restrictions and roles to isolate different teams and the resources available inside a project.




Figure 4: ArgoCD Projects